



Guy Lohman
(<https://wp.sigmod.org/?author=20>)

APRIL 10, 2014

IS QUERY OPTIMIZATION A “SOLVED” PROBLEM?

≡ Databases (<https://wp.sigmod.org/?cat=16>)

Is Query Optimization a “solved” problem? If not, are we attacking should we identify the “right” problems to solve?

I asked these same questions almost exactly 25 years ago, in an ex Workshop on Database Query Optimization that was organized by t Graefe at the Oregon Graduate Center [**Grae 89a**]. Remarkably an of the issues and critical unsolved problems I identified in that brief Researchers continue to attack the wrong problems, IMHO: **they al can**, i.e., that they have ideas for, rather than the ones that **they s** critical to successfully modeling the true cost of plans and choosing importantly, that will avoid choosing a disastrous plan! At the risk o to re-visit these issues, because I’m disappointed that few in the re taken up my earlier challenge.

(<http://wp.sigmod.org/wp-content/uploads/2014/04/divid>

The root of all evil, the Achilles Heel of query optimization, is the es intermediate results, known as cardinalities. Everything in cost estim many rows will be processed, so the entire cost model is predicated model. In my experience, the cost model may introduce errors of al cardinality, but the cardinality model can quite easily introduce erro **magnitude**! I’ll give a real-world example in a moment. With such “Why did the optimizer pick a bad plan?” Rather, the wonder is “Wh pick a decent plan?”

“Well,” you say, “we’ve seen lots of improvements in histograms, w statistics since 1989. Surely we do better now.” There’s been no sh true, but the wealth of such papers precisely illustrates my point. D that improve selectivity estimation for individual local predicates of 47 AND 63” by a few percent doesn’t really matter, when other, mu introduced elsewhere in cardinality estimation dwarf those minor ir engineering, folks. If I have to review one more such paper on impr predicates, I’ll scream (and reject it)! It ***just doesn’t matter!*** Wha enough.

What still introduces the most error in cardinality estimation is (a) t parameter markers, (b) the selectivity of join predicates, and, even how we combine selectivities to estimate the cardinality. Amazingly have enjoyed the least research attention, or at least the fewest nu to solve them, unless I’ve missed some major contributions lately. I topics in turn, describing the fundamental causes of errors, and why reach disastrous proportions, illustrated by war stories from real cu

Host variables and parameter markers

Host variables, parameter markers, and special registers occur in S applications on top of the DBMS, not humans, invoke most queries, papers. Such applications typically get the constants used in predica blank” field on a web page, for example. The SQL predicate then lo :hv1 AND :hv2”. At compile time, the optimizer has no clue what :h cannot look them up in those wonderful histograms that we all have make a wild guess on the average or likely values, which could be c execution to the next, or even due to skew. A war story illustrates t

One of our major ISVs retrofitted a table with a field that identified came from. It had 6 distinct values, but 99.99% of the rows had th subsystem, i.e., when there was only one. A predicate on this subsy to every query, with the value being passed as a host variable. Not priori, DB2’s optimizer used the average value of $1/|\text{distinct values}|$ predicate’s true selectivity was usually 0.9999 (not selective at all) was 0.0001 (extremely selective).

Version 5 of DB2 for OS/390 (shipped June 1997) developed a practical optimization for host variables, parameter markers, and special register bind options REOPT(ALWAYS) and REOPT(ONCE). The latter re-optimizes the statement is executed with actual values for the parameters, and the former will be "typical", whereas the former forces re-optimization every time the statement is run. Later (<http://publib.boulder.ibm.com/infocenter/dzichelp/topic=2Fcom.ibm.db2z10.doc.comref%2Fsrc%2Ftpc%2Fd>) a REOPT(AUTO) option was added to autonomically determine if re-optimization is based upon the change in the estimated filter factors from the last execution.

The paucity of innovation in calculating join predicate selectivities is extant systems still use the techniques pioneered by System R for calculating a join as the inverse of the maximum of the two join-column cardinalities. This essentially assumes that the domain of one column is a subset of the domain of the other. This assumption is valid for key domains having referential integrity constraints. However, the overlap of the two sets may vary greatly depending upon the semantics of the domains. Furthermore, the common practice of pre-populating a dimension table with 100 years of dates into the future can incorrectly bias the selectivity. If a fact table initially has only a few months of data. Statistics on join predicates can give much more precise selectivities for join predicates, if we were to maintain them. The costs of maintenance and lock contention of updates to these join predicate statistics would not solve the **intersection problem** typical of star schemas.

For example, suppose a fact table of Transactions has dimensions for Products and Dates of each transaction. Though current methods provide access for predicates local to each dimension, e.g., `ProductName = 'Docker Jose'` and `Date = '23-Feb-2013'`, it is impossible to determine the effect of the interaction of these predicates on the fact table. Perhaps the San Jose store had a sale of Dockers that day that expired the next day, and a similar sale on some other day, so that the individual selectivities for each day, store, and product are not independent, so that the actual sales of Dockers on the two days would be significantly different. The **interaction** of these predicates, through join predicates in this case, is not addressed. This leads naturally to the final and most challenging problem.

Correlation of columns

With few exceptions ([Chri 83], [Vand 86]), query optimizers since the 1980s have assumed that the selectivities of predicates are independent. These individual selectivities are combined together. Effectively, this assumes that $\text{Prob}\{\text{ColX} = \text{Value1} \text{ and } \text{ColY} = \text{Value2}\} = \text{Prob}\{\text{ColX} = \text{Value1}\} * \text{Prob}\{\text{ColY} = \text{Value2}\}$, i.e., that the columns are probabilistically independent, if you recall your college probability. In the database industry, this assumption is often valid. **However, occasionally**

My favorite example, which occurred in a customer database, is `Make = 'Honda' and Accord = 'Accord'`. To simplify somewhat, suppose there are 10 Makes and 100 Accords. The independence (and uniformity) assumption gives us a selectivity of 0.1. In reality, since only Honda makes Accords, by trademark law, the real selectivity is 1.0. This is a *under-estimate the cardinality by an order of magnitude*. Such optimization errors are often worse than pessimistic over-estimation errors, because they cause certain operations to be cheaper than they really are, causing nasty performance problems. The only way to avoid such errors is for the database administrator to know the semantic relationship (a functional dependency, in this case) between the columns, and to collect **column group statistics**, as **DB2 products now allow** (<http://www.ibm.com/developerworks/data/library/technicalArticles/0612kapoor/>).

To identify these landmines in the schema automatically, Stillger et al. developed the LEarning Optimizer (LEO), which opportunistically and automatically compares actual cardinalities to optimizer estimates, to identify column combination correlation errors. Ilyas et al. [Ilya 04] attacked the problem more

(CORrelation Detection by Sampling), searching somewhat exhaust between any two columns in samples from the data before running and colleagues [Mark 05], [Mark 07] have made ground-breaking way to combine the selectivities of conjuncts in partial results.

All great progress on this problem, but ***none yet solves the problem predicates*** that can be inadvertently introduced by the query write that “more is better”, that providing more predicates helps the DBM American as Apple Pie! Let me illustrate with one of my favorite wa

At a meeting of the International DB2 User’s Group, a chief database major U.S. insurance company whom I’d helped with occasional bac conduct a class on-site. I suggested it include an exercise on a real After my class, she obliged me by presenting two 1-inch stacks of p EXPLAIN of a plan for a query. I feared I was going to embarrass m under the gun. The queries differed in only one predicate, she said, seconds whereas the one with the extra predicate took over an hou

I instinctively examined first the cardinality estimates for the result slower one had a cardinality estimate 7 orders of magnitude less th asked what column the extra predicate was on, my host explained t key constructed of the first four letters of the policy-holder’s last na initials, the zip code, and last four digits of his/her Social Security N query have predicates on all those columns? Of course! And how m table? Ten million. Bingo! I explained that that predicate was compl others, and its selectivity, $1/10^7$, when multiplied by the others, un cardinality by 7 orders of magnitude, wreaking havoc with the plan. minutes to figure this out, and I was immediately dubbed a “genius straightforward: the added predicate might help the run-time, espe on that column, but it totally threw off the optimizer, which couldn’t that redundancy without LEO or CORDS.

So c’mon, folks, let’s attack problems that really matter, tho optimizer disasters, and stop polishing the round ball.

Disclaimer: The postings on this site are my own and don’t necessa positions, strategies or opinions.

References

- [Chri 83] S. Christodoulakis, “Estimating record selectivities”, Info. Systems 8,2 (1983).
- [Grae 89a] G. Graefe, editor, Workshop on Database Query Optimization, CSE Tec Graduate Center, Portland, OR, 30 May 1989.
- [Grae 89b] G. Graefe, “The stability of query evaluation plans and dynamic query of ACM-SIGMOD, Portland, OR (1989).
- [Ioan 97] Y. Ioannidis et al., “Parametric query optimization”, VLDB Journal, 6(2), (1997).
- [Ilya 04] I. F. Ilyas, V. Markl, P. J. Haas, P. Brown, A. Aboulnaga: CORDS: Automatic Soft Functional Dependencies. SIGMOD Conference 2004: 647-658.
- [Mark 05] V. Markl, N. Megiddo, M. Kutsch, T. Minh Tran, P. J. Haas, U. Srivastava: Selectivity of Conjunctions of Predicates. VLDB 2005: 373-384.
- [Mark 07] V. Markl, P. J. Haas, M. Kutsch, N. Megiddo, U. Srivastava, T. Minh Tran: Cost estimation via maximum entropy. VLDB J. 16(1): 55-76 (2007)
- [Redd 05] N. Reddy and J. R. Haritsa. “Analyzing plan diagrams of database query optimization”, Proc. of the 31st ACM SIGMOD Conference on Management of Data, 2005.
- [Seli 79] P.G. Selinger, M.M. Astrahan, D.D. Chamberlin, R.A. Lorie, and T.G. Price, ‘Relational Database Management System”, Procs. Of ACM-SIGMOD (1979), pp. 23-32.
- [Stil 01] M. Stillger, G. M. Lohman, V. Markl, M. Kandil: LEO – DB2’s LEarning Optimal Query Plan. Technical report (http://www.cs.columbia.edu/~wfan/PAPERS/Leo.pdf).
- [Stoy 08] J. Stoyanovich, K. A. Ross, J. Rao, W. Fan, V. Markl, G. Lohman: ReoptSMA Cache. **Technical report** (<http://www.cs.columbia.edu/~wfan/PAPERS/ReoptSMA.pdf>).
- [Vald 87] P. Valuriez, “Join indices”, ACM Trans. on Database Systems 12, 2 (June 1987).
- [Vand 86] B.T. Vander Zanden, H.M. Taylor, and D. Bitton, “Estimating block access correlated”, Procs. of the 12th Intl. Conf. on Very Large Data Bases (Kyoto, Sept. 1986).

Blogger’s Profile:

Dr. Guy M. Lohman (<http://researcher.ibm.com/researcher/view.php?person=DrGuyM.Lohman>) is a Senior Research Manager in the Disruptive Information Management Architectures in the Advanced Information Management Research Division’s Almaden Research Center in San Jose, California, where he has currently manages the Blink research project, which contributed **BLU Acceleration** (<http://www.ibm.com/software/data/db2/linux-unix-windows/db2-blu-acceleration>) and Windows (LUW) 10.5 (GA’d 2013) and the query engine of the IBM Smart Analytics V1.1 and the **Informix Warehouse Accelerator** (<http://www.ibm.com/software/data/infomix/advanced-enterprise-edition>) products (2007-2010). Dr. Lohman was the architect of DB2 LUW and was responsible for its development from 1992 to 1997 (versions 2 - 9.7), the prototyping of Visual Explain, efficient sampling, the DB2 Index Advisor, and optimizer for DB2. Dr. Lohman was elected to the **IBM Academy of Technology** (<http://www.ibm.com/ibm/academy/index.html>) in 2002 and made an IBM Most Valuable Professional the General Chair for **ACM’s 2013 Symposium on Cloud Computing** (<http://www.acmcloud.org>) and General Co-Chair of the **2015 IEEE International Conference on Data Engineering** (<http://www.icde2015.kr>). Previously, he was on the editorial boards of the “Very Large Databases”, “Distributed and Parallel Databases”. He is the author of over 75 papers in the refereed journals and has been awarded 39 U.S. patents. His current research interests involve disruptive Business Intelligence, advanced data analytics, query optimization, self-managing database management appliances, database compression, and autonomic problem determination.

Copyright © 2014, Guy M. Lohman, All rights reserved.



1,556 views

Systems & Databases: Let’s Break Down the Walls (<https://wp.sigmod.org/?p=1009>)

Exploratory search: New name for an old hat? (<https://wp.sigmod.org/?p=1183>)

8 Comments



James Parker on April 17, 2014

I found this an interesting and useful paper, given that this is not my area of expertise (main-memory DBMS). However, I wonder to what extent query optimization ought to also consider the working set of rows already in memory when determining a query plan as well. Such an approach might even improve performance when paging of the underlying in-VM rowss is possible, by keeping them in the working set(s) of the DBMS process(es).




Stephen Dillon (<http://www.linkedin.com/in/stephendillon/>) on April 20, 2014

James,

I too was contemplating the role of an in-memory DBMS (IMDBS) and these query optimization points. Just how much of an impact will they contribute to a system such as VoltDB? This is actually something I've had discussions with other architects about to address their own proprietary main memory DBMS.

My initial thoughts are that a commercial IMDBS may be impacted less, than a traditional RDBMS, by such optimizations when you consider the significant impact main-memory operations already provide queries. By commercial, I really mean one that is professionally developed and properly addresses latching, locking, durability, et al. That does not mean proper or better optimization is not needed for an IMDBS. Nobody has proven these optimizations yet with an IMDBS to my knowledge. I just believe that as compared to a traditional RDBMS, an IMDBS may not benefit as much. I also believe that these modern main-memory DBs such as memSQL or VoltDB allow “some” bad code to hide under the guise of fast execution plans that have benefited from zero disk reads. Now for someone's proprietary IMDBS well these optimizations could certainly provide more benefits considering how less efficient it is to a commercial product.



Ruslan Fomkin (<http://se.linkedin.com/in/fomkin/>) on April 25, 2014

I haven't tested, but in my opinion the query optimization has similar effect on in-memory and disk-based databases. I see two reason for this:

- (1) Data in disk-based database are cached in main-memory. Thus no disk access in both DBMS types. The main difference is cardinality units: pages vs cache lines.
- (2) The main bottleneck for query execution is memory wall nowadays. Thus suboptimal query plans in IMDBMS can still perform by an order of magnitude worse then optimal one.



James Parker on April 26, 2014

Stephen,

I think you misunderstood me; I was not thinking directly of main memory (AKA in-memory) DBMSs, but rather about the memory hierarchy (memory vs. “disk”) and where specific data resides at the time of a query. In a traditional DBMS, some data is known to be in memory, including which rows from which tables — a query optimizer might favor joins using tables with most or all rows in memory over those which are not, as a part of query optimization. Of course the effort must take into account paging, if it may occur; however paging can often be managed programmatically, e.g., by “locking” pages in memory.

The similar issue in a main memory DBMS would be where in the CPU cache hierarchy data resides; this is generally of much lower utility, since machine architectures generally do not expose ways to manage these caches programmatically (although I have speculated as to how hardware and OS architects might provide useful features toward this end, and would almost certainly need to be combined with processor scheduling and interrupt handling).



spaghettidba (<http://spaghettidba.com>) on April 18, 2014

Thanks for the interesting writing.

The recently released SQL Server 2014 incorporates a new cardinality estimator that (partially) addresses the predicate independence issue. The new algorithm used is known as “exponential backoff” and it helps reducing the estimation error with multiple predicates. You can read about it here: <http://www.sqlperformance.com/2014/01/sql-plan/cardinality-estimation-for-multiple-predicates> (<http://www.sqlperformance.com/2014/01/sql-plan/cardinality-estimation-for-multiple-predicates>)



Lothar Flatz on April 18, 2014

After 15 years of tuning in the field I have come across all of these issues.

This I have my own examples i.e. Correlation of columns: mobile phone service provider. Primary key of table units consist of customer number and phone number. Selectivity of join predicates: Road assistance search for service centers. Selectivity differs largely depending on country of accident. Wrote a paper on it. Hope it gets accepted. 😊



Alexandr Savinov (<http://conceptoriented.org/savinov>) on April 18, 2014

Thanks for the interesting post and relevant works. One problem in query optimization I am currently dealing with is optimizing column operations for analytical queries. It is highly important for a novel system for analytical data integration (ConceptMix – <http://conceptoriented.com> (<http://conceptoriented.com>)) where all operations with data are described in terms of an algebra of functions (a function represents a column).

Essentially, all manipulations are reduced to manipulating functions (and inverse functions) and the problem is to translate queries from the concept-oriented expression language to these functional operations.



Radim Baca on April 22, 2014

Thank you for an inspirational article. I have a question regarding the “redundant predicates problem”: can the CORDS method solve the problem? From my perspective this problem is just a result of high correlation of attributes, which is addressed by CORDS. Thanks

Comments are closed

Search

Categories

Select Category



Recent Comments

Atsuyuki Morishima on **Imagine all the People and AI in the Future of Work**
(<https://wp.sigmod.org/?p=2931#comment-54638>)

Zakaria Maamar on **Imagine all the People and AI in the Future of Work**
(<https://wp.sigmod.org/?p=2931#comment-54635>)

Yuriy on The Rise of Natural Language Interfaces to Databases (<https://wp.sigmod.org/?p=2897#comment-54545>)

william e winkler (<https://www.census.gov/srd/csrreports/byyear.html>) on Data cleaning is a machine learning problem that needs data systems help! (<https://wp.sigmod.org/?p=2288#comment-54417>)

Gio Wiederhold (<http://i.stanford.edu/~gio>) on A Tribute to José Alfredo Blakeley (<https://wp.sigmod.org/?p=2261#comment-54399>)

Archives

September 2019 (<https://wp.sigmod.org/?m=201909>)

June 2019 (<https://wp.sigmod.org/?m=201906>)

November 2018 (<https://wp.sigmod.org/?m=201811>)

October 2018 (<https://wp.sigmod.org/?m=201810>)

August 2018 (<https://wp.sigmod.org/?m=201808>)

June 2018 (<https://wp.sigmod.org/?m=201806>)

April 2018 (<https://wp.sigmod.org/?m=201804>)

March 2018 (<https://wp.sigmod.org/?m=201803>)

February 2018 (<https://wp.sigmod.org/?m=201802>)

June 2017 (<https://wp.sigmod.org/?m=201706>)

May 2017 (<https://wp.sigmod.org/?m=201705>)

April 2017 (<https://wp.sigmod.org/?m=201704>)

February 2017 (<https://wp.sigmod.org/?m=201702>)

January 2017 (<https://wp.sigmod.org/?m=201701>)

October 2016 (<https://wp.sigmod.org/?m=201610>)

April 2016 (<https://wp.sigmod.org/?m=201604>)
February 2016 (<https://wp.sigmod.org/?m=201602>)
November 2015 (<https://wp.sigmod.org/?m=201511>)
September 2015 (<https://wp.sigmod.org/?m=201509>)
July 2015 (<https://wp.sigmod.org/?m=201507>)
June 2015 (<https://wp.sigmod.org/?m=201506>)
March 2015 (<https://wp.sigmod.org/?m=201503>)
February 2015 (<https://wp.sigmod.org/?m=201502>)
December 2014 (<https://wp.sigmod.org/?m=201412>)
November 2014 (<https://wp.sigmod.org/?m=201411>)
August 2014 (<https://wp.sigmod.org/?m=201408>)
June 2014 (<https://wp.sigmod.org/?m=201406>)
April 2014 (<https://wp.sigmod.org/?m=201404>)
February 2014 (<https://wp.sigmod.org/?m=201402>)
December 2013 (<https://wp.sigmod.org/?m=201312>)
October 2013 (<https://wp.sigmod.org/?m=201310>)
September 2013 (<https://wp.sigmod.org/?m=201309>)
June 2013 (<https://wp.sigmod.org/?m=201306>)
May 2013 (<https://wp.sigmod.org/?m=201305>)
March 2013 (<https://wp.sigmod.org/?m=201303>)
December 2012 (<https://wp.sigmod.org/?m=201212>)
November 2012 (<https://wp.sigmod.org/?m=201211>)
August 2012 (<https://wp.sigmod.org/?m=201208>)

July 2012 (<https://wp.sigmod.org/?m=201207>)

May 2012 (<https://wp.sigmod.org/?m=201205>)

April 2012 (<https://wp.sigmod.org/?m=201204>)

February 2012 (<https://wp.sigmod.org/?m=201202>)

Blog editor: Georgia Koutrika

ACM SIGMOD Blog

Powered by **WordPress** (<http://wordpress.org/>) | Theme by **WpFreeware** (<https://www.wpfreeware.com/>)