

Pivotal.

# The Internals of GPORCA Optimizer

---

Xin Zhang (xzhang@pivotal.io)

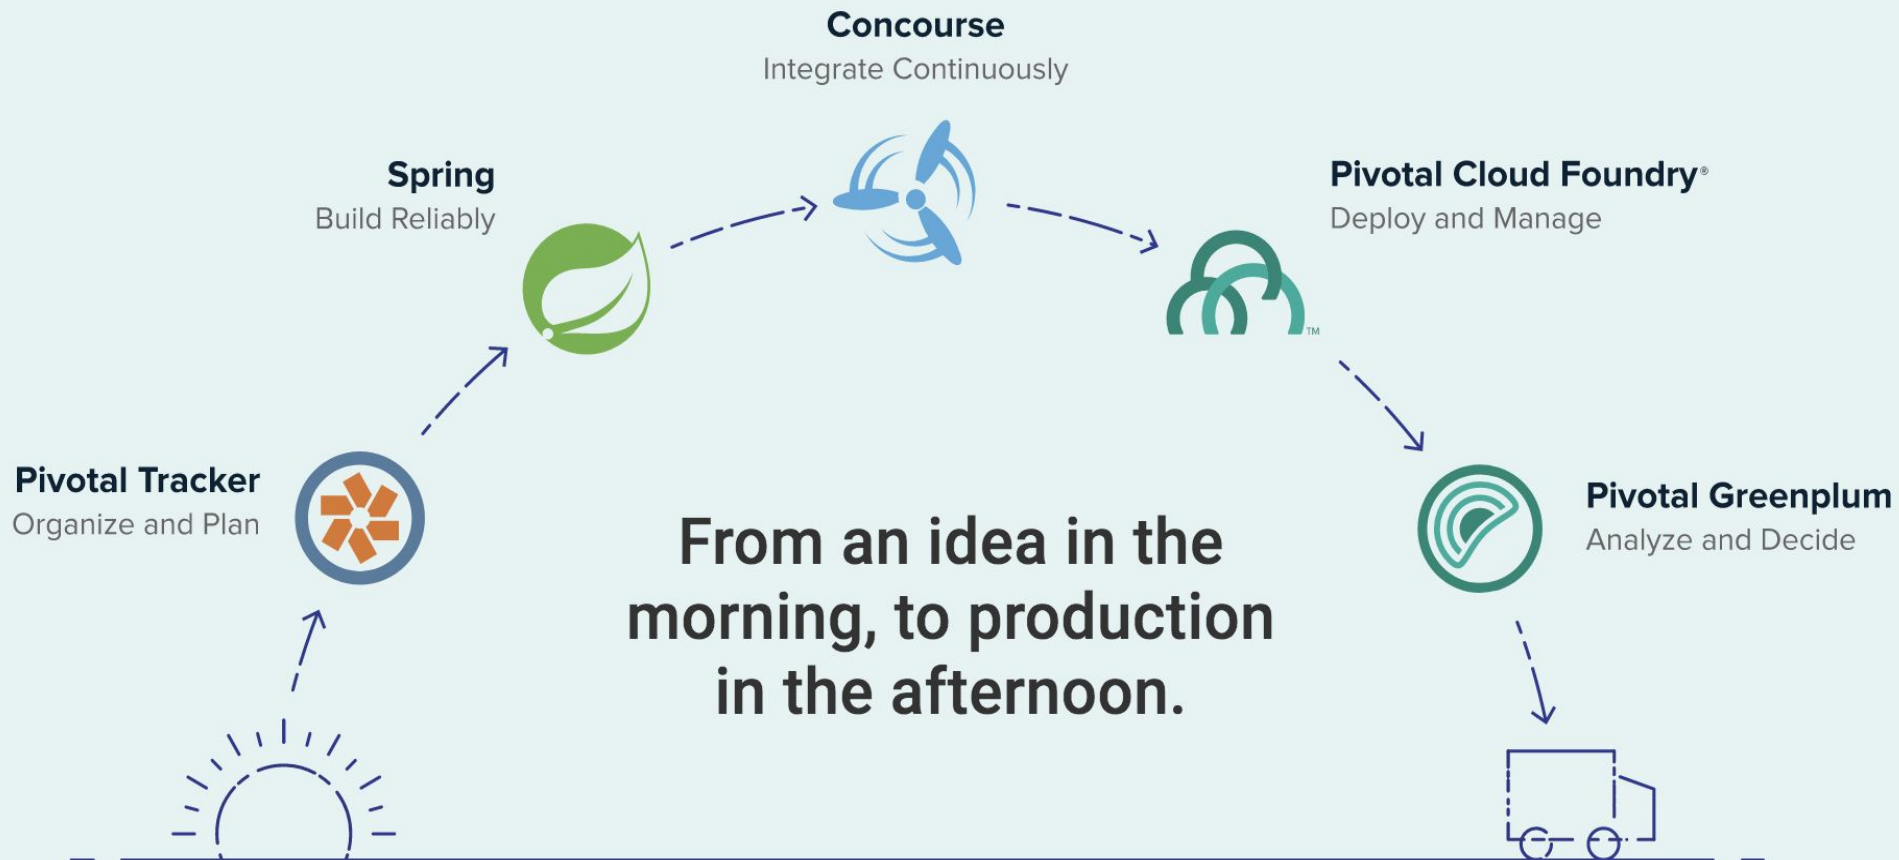
PGConf Seattle 2017  
Nov 2017



# Disclaimer

This presentation contains statements relating to Pivotal's expectations, projections, beliefs and prospects which are "forward-looking statements" about Pivotal's future which by their nature are uncertain. Such forward-looking statements are not guarantees of future performance, and you are cautioned not to place undue reliance on these forward-looking statements. Actual results could differ materially from those projected in the forward-looking statements as a result of many factors, including but not limited to: (i) adverse changes in general economic or market conditions; (ii) delays or reductions in information technology spending; (iii) risks associated with managing the growth of Pivotal's business, including operating costs; (iv) changes to Pivotal's software business model; (v) competitive factors, including pricing pressures and new product introductions; (vi) Pivotal's customers' ability to transition to new products and computing strategies such as cloud computing, the uncertainty of customer acceptance of emerging technologies, and rapid technological and market changes; (vii) Pivotal's ability to protect its proprietary technology; (viii) Pivotal's ability to attract and retain highly qualified employees; (ix) Pivotal's ability to execute on its plans and strategy; and (x) risks related to data and information security vulnerabilities. All information set forth in this presentation is current as of the date of this presentation. These forward-looking statements are based on current expectations and are subject to uncertainties and changes in condition, significance, value and effect as well as other risks disclosed previously and from time to time in documents filed by Dell Technologies Inc., the parent company of Pivotal, with the U.S. Securities and Exchange Commission. Dell and Pivotal assume no obligation to, and do not currently intend to, update any such forward-looking statements after the date of this presentation.

The following is intended to outline the general direction of Pivotal's offerings. It is intended for information purposes only and may not be incorporated into any contract. Any information regarding pre-release of Pivotal offerings, future updates or other planned modifications is subject to ongoing evaluation by Pivotal and is subject to change. This information is provided without warranty or any kind, express or implied, and is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions regarding Pivotal's offerings. These purchasing decisions should only be based on features currently available. The development, release, and timing of any features or functionality described for Pivotal's offerings in this presentation remain at the sole discretion of Pivotal. Pivotal has no obligation to update forward-looking information in this presentation.



# Open Source Software (OSS)



**Oct 2015**, Greenplum Database Open Sourced

(based on **PostgreSQL 8.2**)

**Sep 7th 2017**, Greenplum Database **5.0** w/ GPORCA

(based on **PostgreSQL 8.3**)

**Oct 23rd 2017**, Greenplum Database **6.0 Alpha**

(based on **PostgreSQL 8.4**)

...

...

Pivotal

**Sep 2015**, Shin joined Pivotal



**Jan 2016**, GPORCA Open Source V1.636

<http://engineering.pivotal.io/post/gporca-open-source/>

**Jan 2017**, GPORCA V2.0 (merge GPOS)

**Sep 6th 2017**, GPORCA V2.42.0

**Nov 8th 2017**, GPORCA V2.48.6



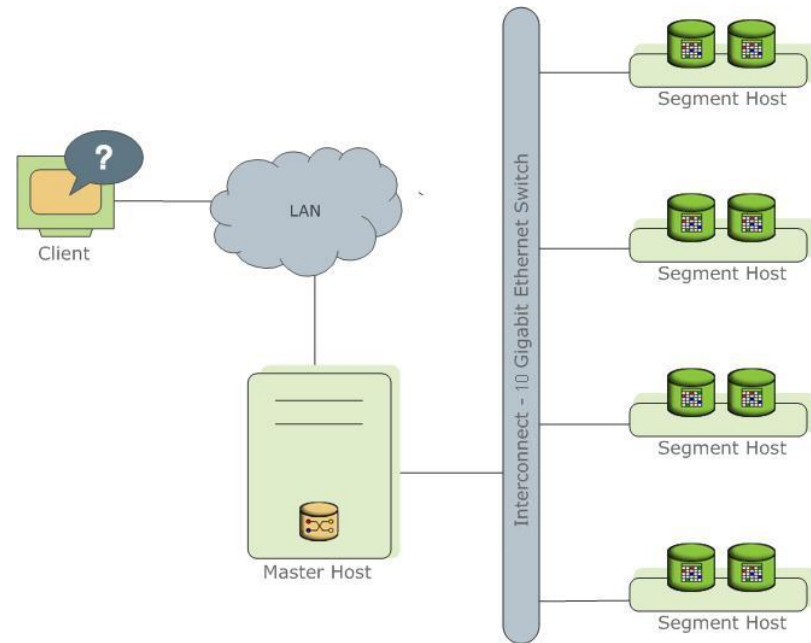
# Greenplum (GPDB): MPP + PostgreSQL

- Shared Nothing Architecture
- Data Distributed on Cluster
- Query Processed Locally in Parallel
- Each Segment is a PostgreSQL Instance



**GREENPLUM  
DATABASE**

<http://greenplum.org/>



# MPP Query Processing Example 1/3



```
SELECT s.beer, s.price  
FROM Bars b, Sells s  
WHERE b.name = s.bar  
AND b.city = 'San Francisco'
```

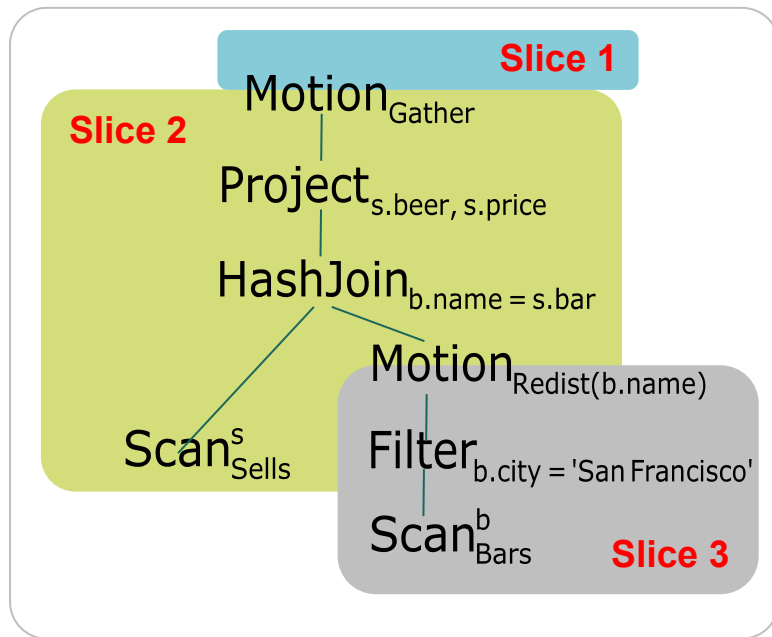
- Bars is distributed randomly → any bar on any host
- Sells is distributed by bar → sells of same bar on same host

# MPP Query Processing Example 2/3

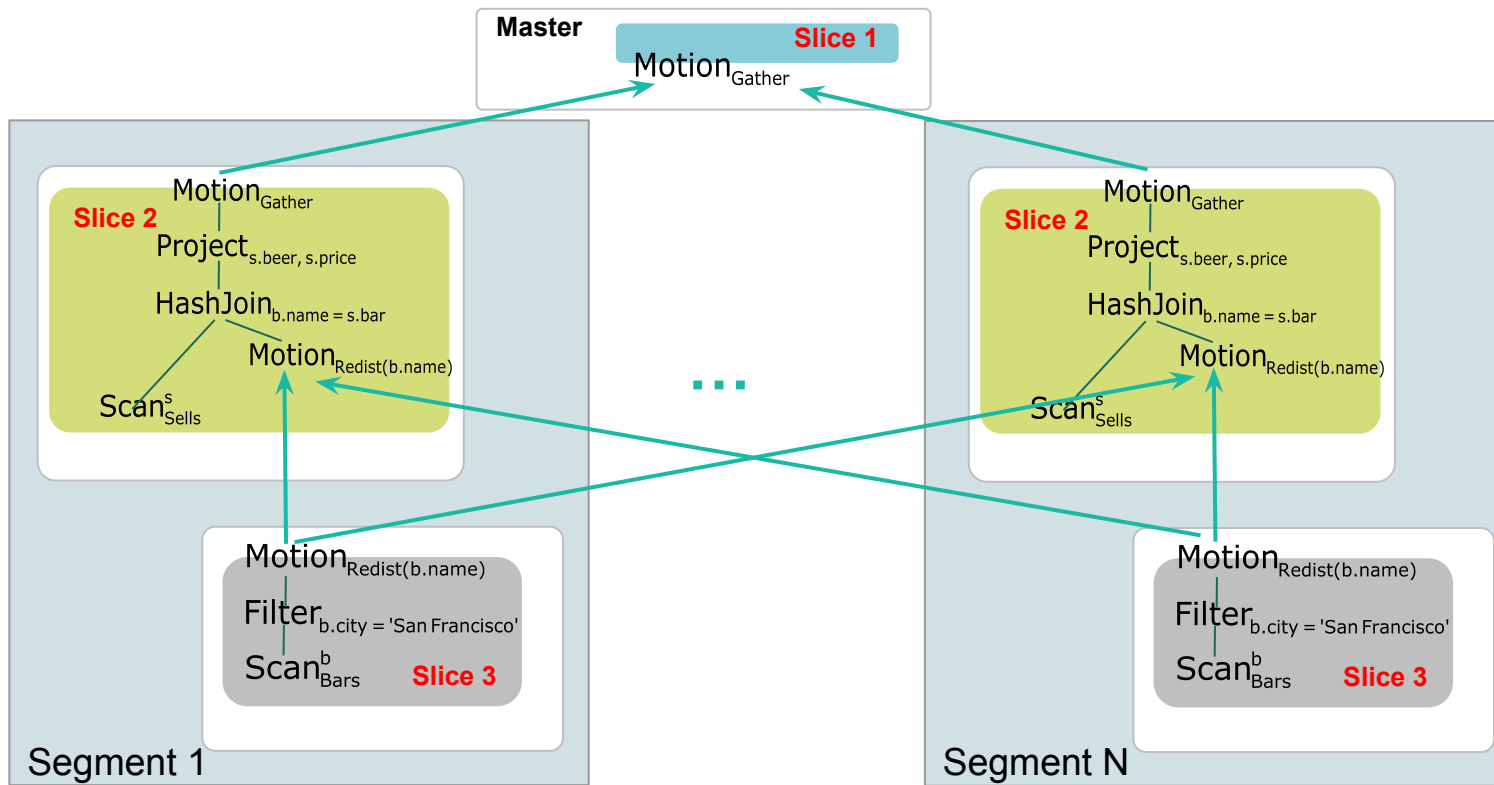


```
SELECT s.beer, s.price
FROM Bars b, Sells s
WHERE b.name = s.bar
AND b.city = 'San Francisco'
```

- Bars is distributed randomly
- Sells is distributed by bar



# MPP Query Processing Example 3/3





A group of people in a workshop setting. A man on the left is pointing at a wall covered in papers. A group of people are sitting on stools in the center, and another man is standing on the right. The image has a dark blue overlay.

# GPORCA find the best MPP plan



# Outline



Context

Architecture

Develop Walkthrough

Build & Test

Contribute



GPORCA



**CONTEXT**

# GPORCA = Greenplum ORCA = ...



Jul 2015, **PQO**: Pivotal Query Optimizer (Released in GPDB 4.3.5.0)

Jun 2014, **ORCA**: SIGMOD 2014 Paper

In source code **GPOPT**

```
src/backend/gpopt
```

```
psql: select gp_opt_version();
```



GPORCA can optimize all

# 99 TPC-DS Queries

(and variations)

Reference: Orca: A Modular Query Optimizer Architecture for Big Data, SIGMOD 2014

# Great SQL Surface of GPORCA



~60 Logical Operators

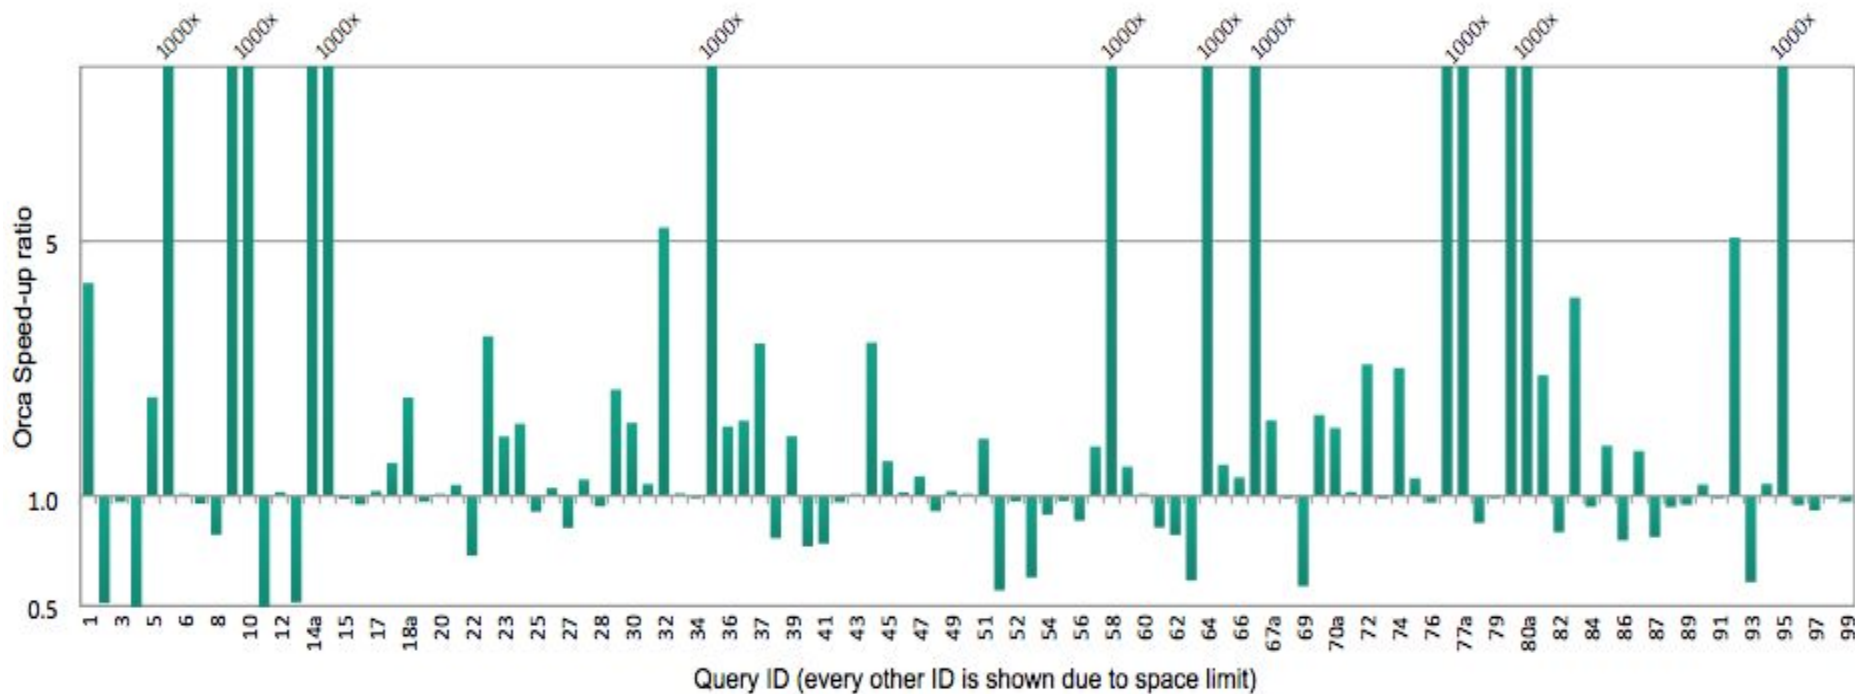
~50 Physical Operators

~40 Scalar Operators

~110 Transformation Rules



# Great Query Performance vs. Planner



TPC-DS 10TB, 16 nodes, 48 GB/node



# Key Features vs. Planner

- Smarter partition elimination
- Multi-level partitioning support (4.3.6)
- Subquery unnesting
- Common table expressions (CTE)
- Join on castable data types
- Improved join ordering (better join order)
- Join-Aggregate reordering (push/pull AGG over join)
- Sort order optimization (avoid sort multiple times)
- Skew awareness (redistribution vs. broadcast)



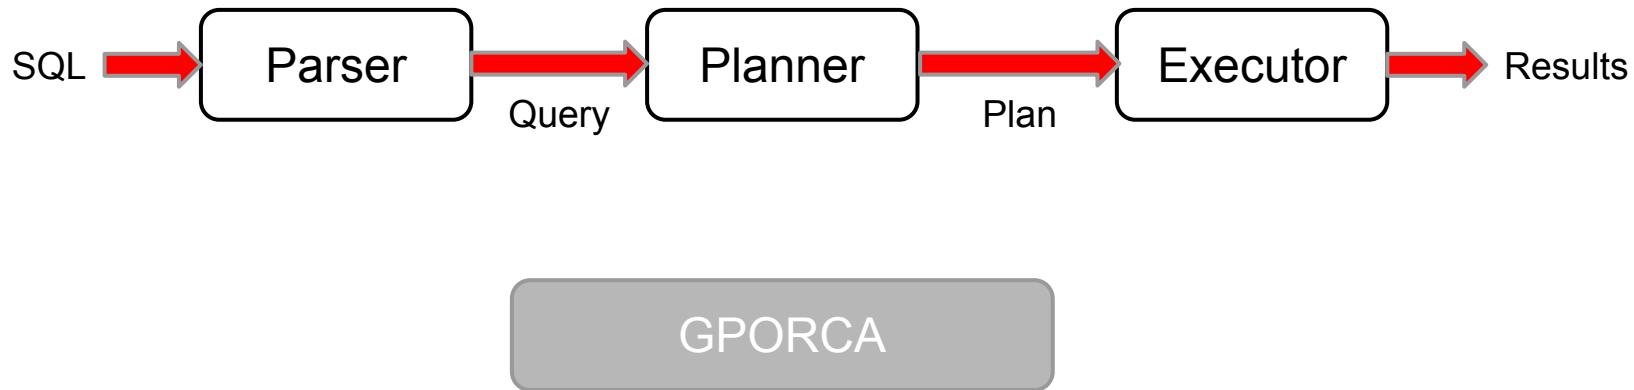


GPORCA

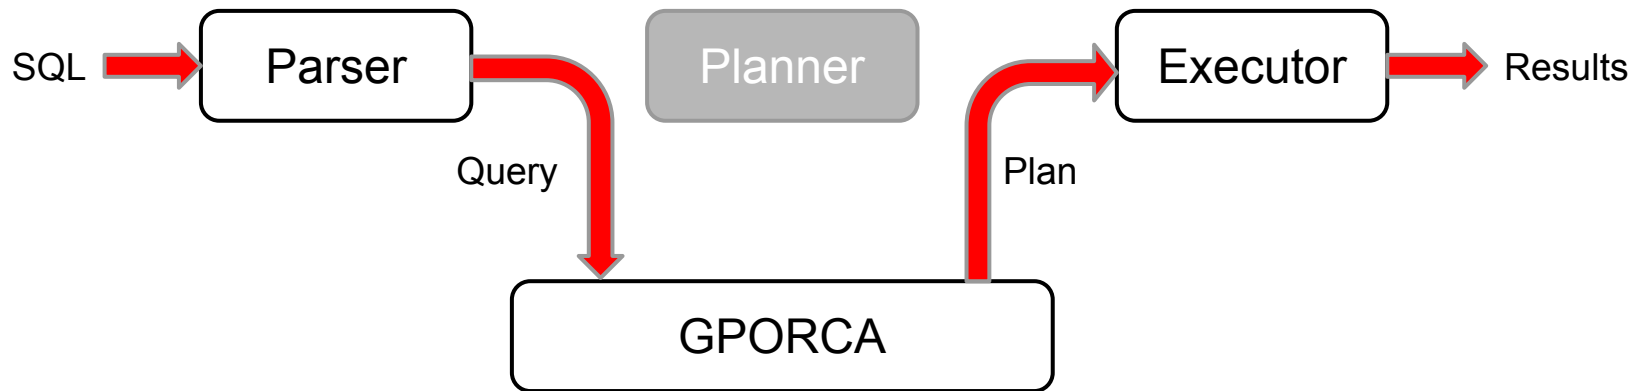
---

# ARCHITECTURE

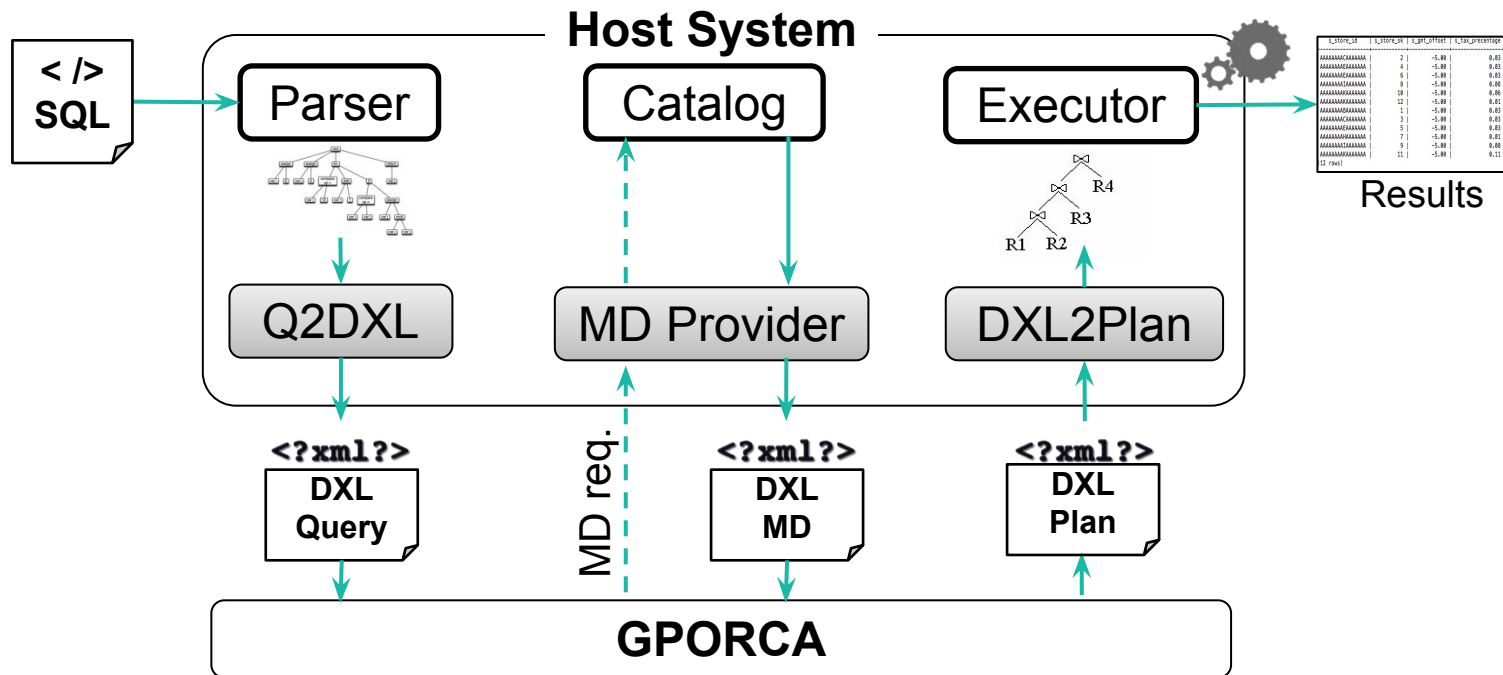
# GPDB Using Planner (`set optimizer=off`)



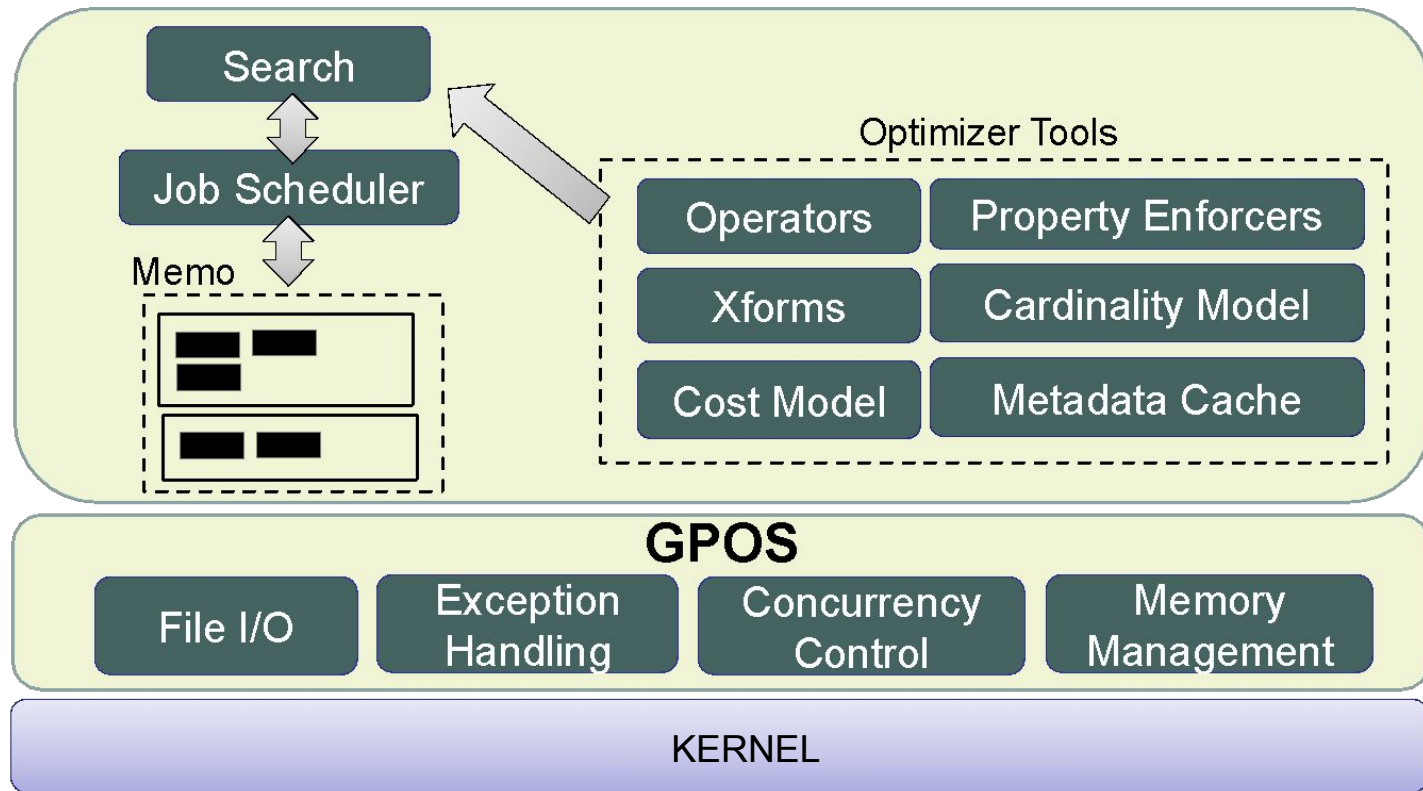
# GPDB Using GPORCA (`set optimizer=on`)



# Multi-Host with DXL (Data eXchange Language) API



# GPORCA Architecture



# Five Optimization Steps in Orca



[Step 0] Pre-Process: e.g., predicates pushdown

[Step 1] Exploration: All equivalent logical plans

[Step 2] Statistics Derivation: histograms

[Step 3] Implementation: Logical to Physical

[Step 4] MPP Optimization: Enforcing and Costing



# Step 0 Pre-Process with 25 iterations (TL, DR)

- (1) Remove unused CTE anchors
- (2) Remove intermediate superfluous limit
- (3) Trim unnecessary existential subqueries
- (4) Remove superfluous outer references from the outer spec in limits, grouping columns, partition/order columns in window operators
- (5) Remove superfluous equality
- (6) Simplify quantified subqueries
- (7) Preliminary unnesting of scalar subqueries
- (8) Unnest AND/OR/NOT predicates and ensure predicates are array IN or NOT IN where applicable
- (9) Infer predicates from constraints**
- (10) Eliminate self comparisons
- (11) Remove duplicate AND/OR children
- (12) Factorize common expressions

- (13) Infer filters out of components of disjunctive filters
  - (14) Pre-process window functions
  - (15) Collapse cascaded union/union all
  - (16) Eliminate unused computed columns
  - (17) Normalize expression
  - (18) Transform outer join into inner join
  - (19) Collapse cascaded inner joins
  - (20) Generate more predicates from constraints**
  - (21) Eliminate empty subtrees
  - (22) Collapse cascade of projects
  - (23) Insert project for scalar subquery return outer reference
  - (24) Reorder children of scalar comparison operator
- Rewrite IN subquery to EXIST subquery with a predicate

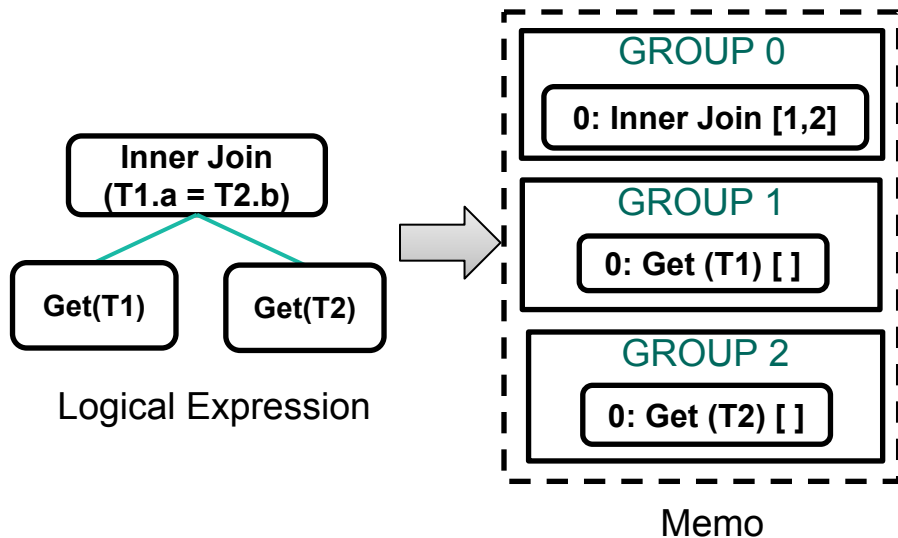
# Step 1 Exploration 1/2: Memoization



Compact in-memory data structure capturing plan space:

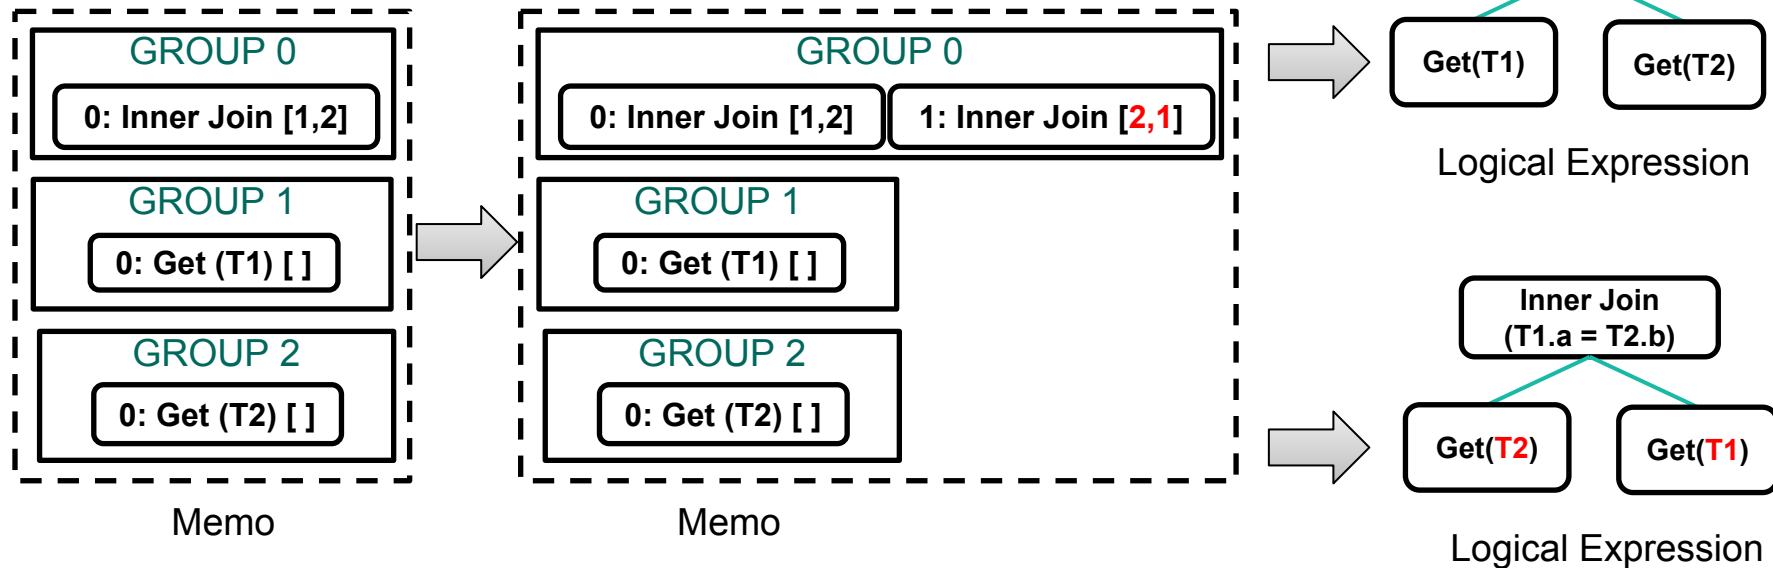
**Group:** Container of equivalent expressions

**Group Expression:** Operator that has other groups as its children

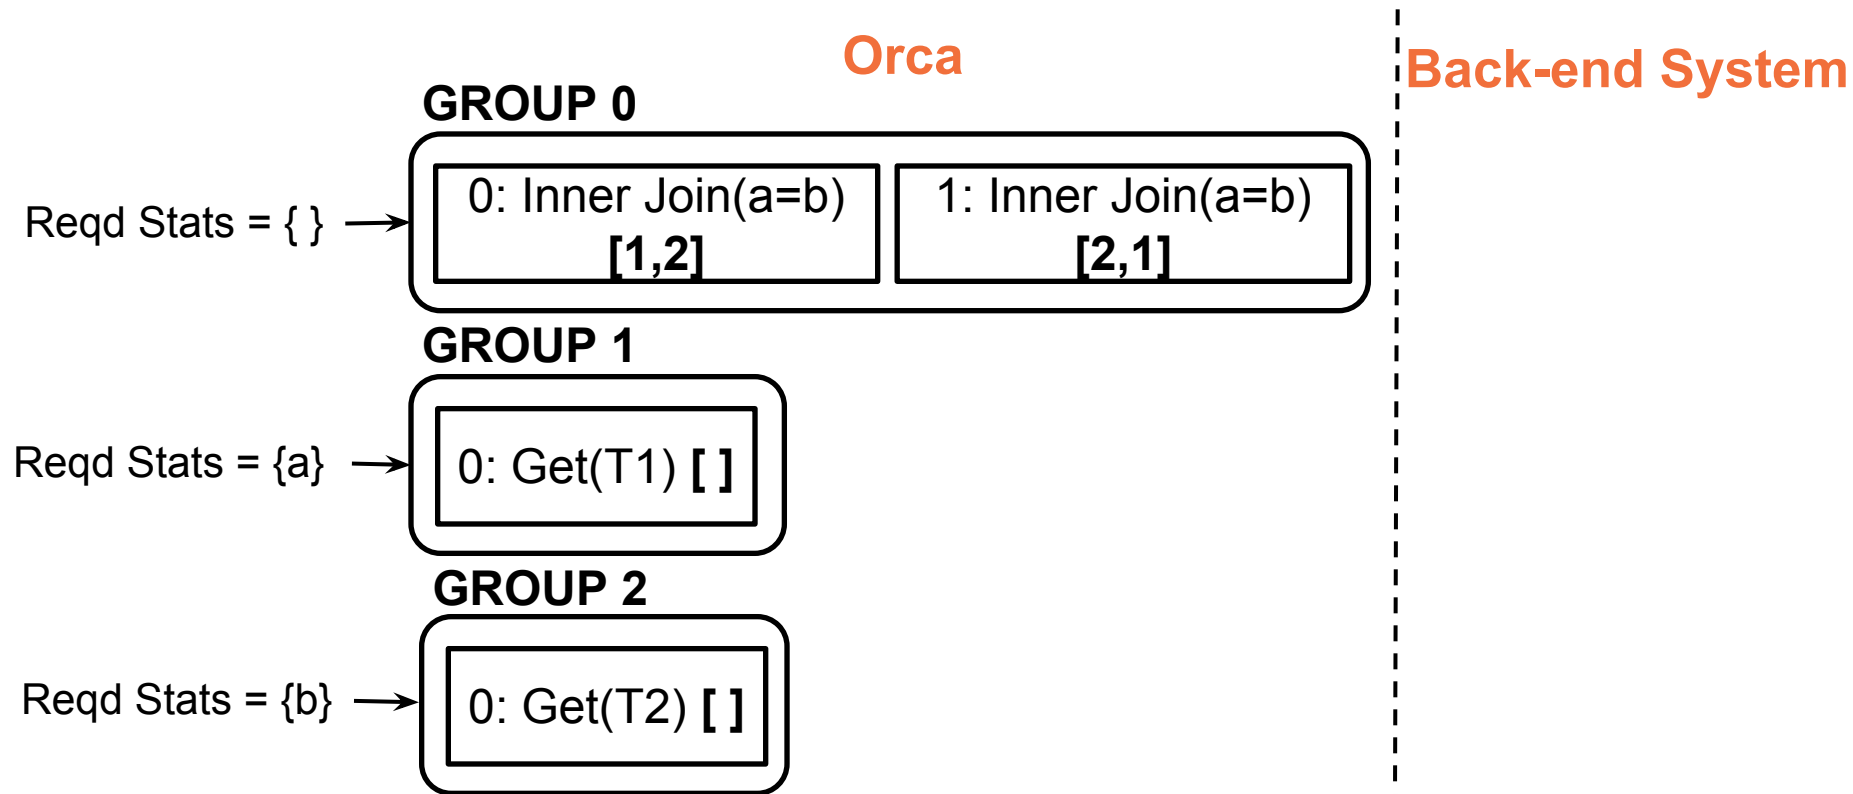




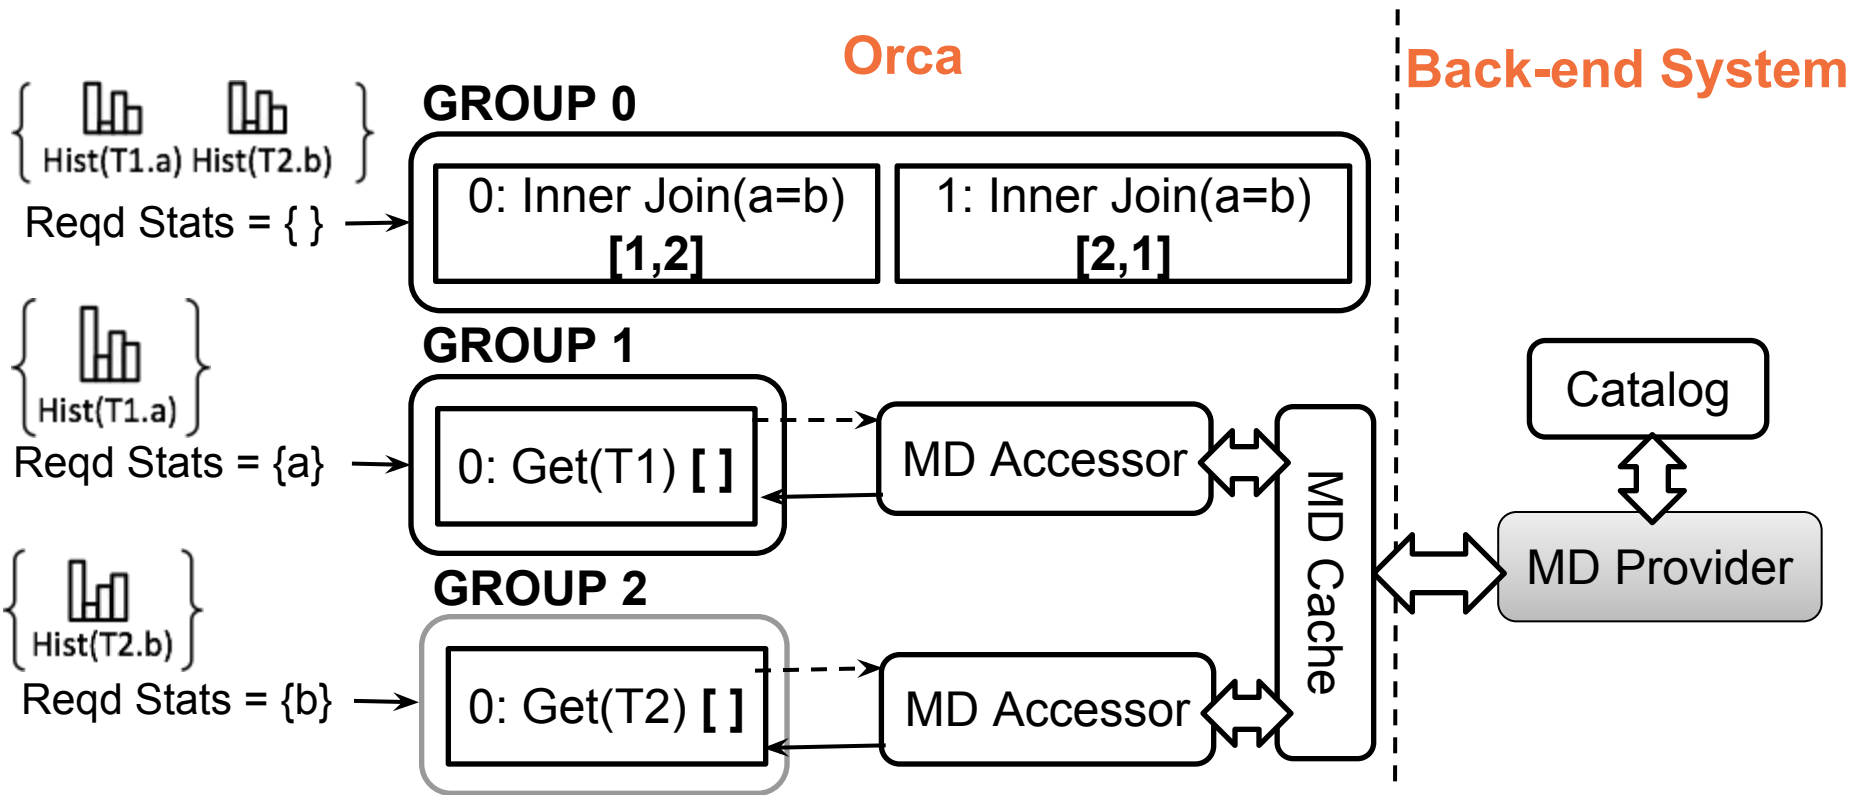
# Step 1 Exploration 2/2: Transformation



# Step 2 Statistics Derivation 1/2



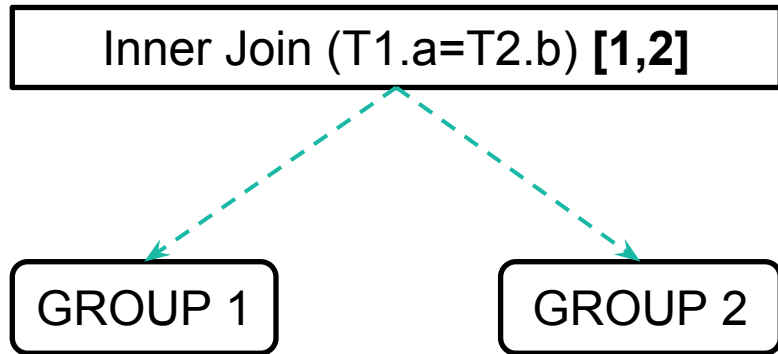
## Step 2 Statistics Derivation 2/2



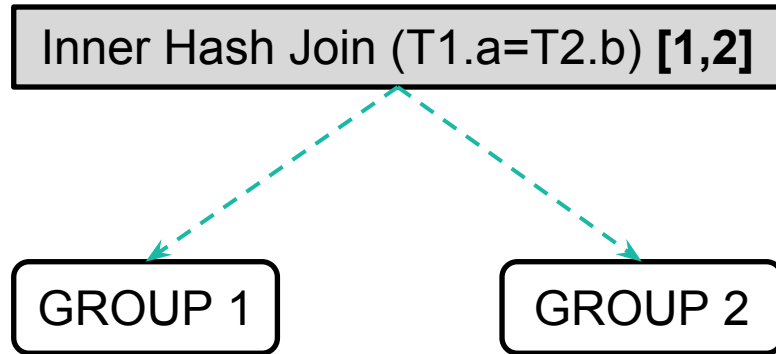
# Step 3 Implementation



Logical Expression



Physical Expression



# Step 4 MPP Optimization 1/3 Requirement



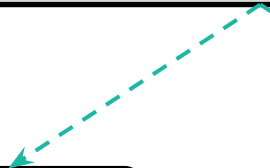
## Optimization Request

Distribution, Order

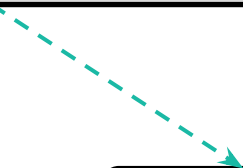
Reqd Props: {Singleton, <T1.a>}



Inner Hash Join (T1.a=T2.b) [1,2]



GROUP 1



GROUP 2

# Step 4 MPP Optimization 2/3 Costing



## Optimization Request

Distribution, Order

Reqd Props: {Singleton, <T1.a>}



Inner Hash Join (T1.a=T2.b) [1,2]



{Hashed(T1.a), Any}

{Hashed(T2.b), Any}

Scan(T1)

**Redistribute(T2.b)**

Scan(T2)

# Step 4 MPP Optimization 3/3 Enforcement



## Optimization Request

Distribution, Order

Reqd Props: {Singleton, <T1.a>}

Inner Hash Join (T1.a=T2.b) [1,2]

Property  
Enforcing

{Hashed(T1.a), Any}

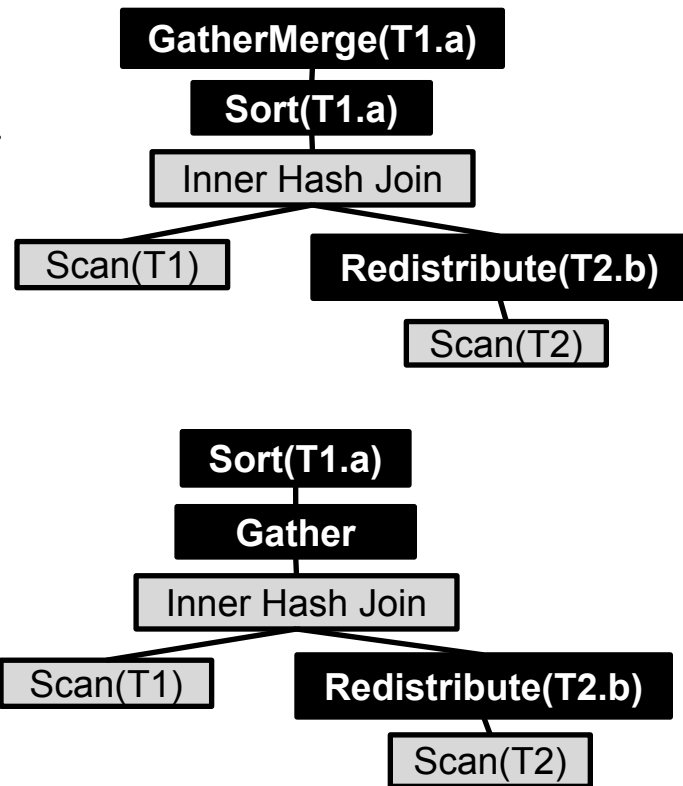
{Hashed(T2.b), Any}

Scan(T1)

Redistribute(T2.b)

Scan(T2)

Pivotal





GPORCA



**DEVELOP**





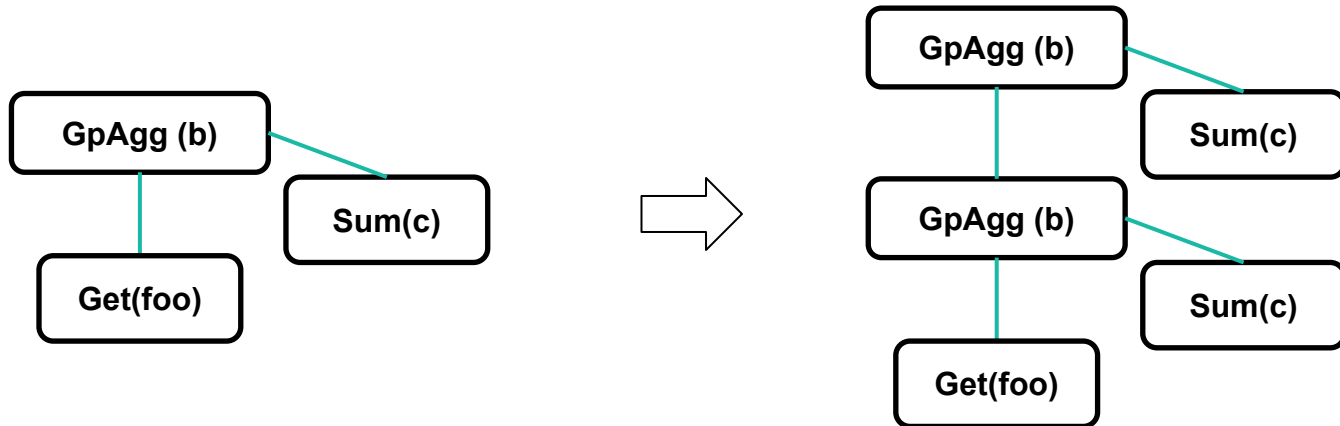
Split an aggregate into a pair of **local** and **global** aggregate.

```
CREATE TABLE foo (a int, b int, c int)
DISTRIBUTED BY (a);
```

```
SELECT sum(c) FROM foo GROUP BY b
```

Do **local** aggregation on segments  
The **global** aggregation on master

# Split Groupby Aggregate



# CXformSplitGbAgg



```
// HEADER FILES
```

```
~/orca/libgpopt/include/gpopt/xforms
```

```
// SOURCE FILES
```

```
~/orca/libgpopt/src/xforms
```

# Transformation Trigger

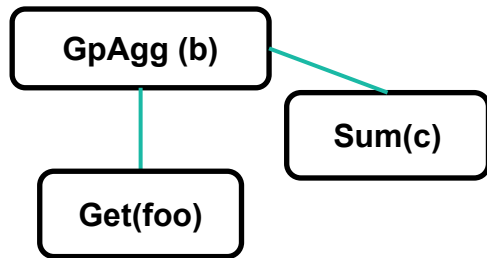


- Pattern
- Pre-Condition Check

# Pattern

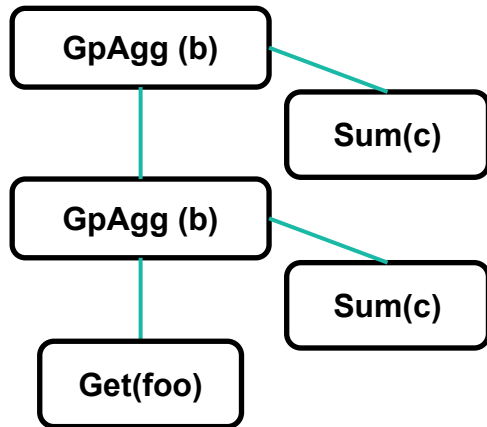


```
GPOS_NEW(pmp)
CExpression
(
    pmp,
    // logical aggregate operator
    GPOS_NEW(pmp) CLogicalGbAgg(pmp),
    // relational child
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternLeaf(pmp)),
    // scalar project list
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternTree(pmp))
));
```



# What's WRONG of this pattern?

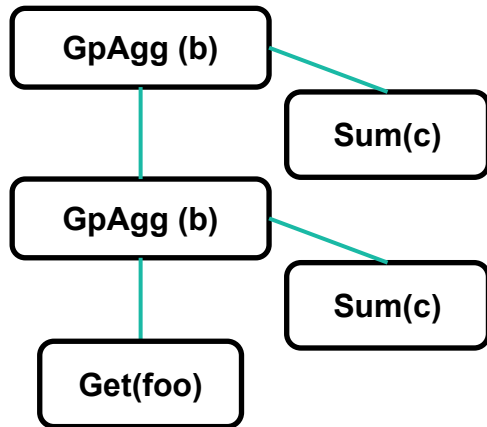
```
GPOS_NEW(pmp)
CExpression
(
    pmp,
    // logical aggregate operator
    GPOS_NEW(pmp) CLogicalGbAgg(pmp),
    // relational child
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternLeaf(pmp)),
    // scalar project list
    GPOS_NEW(pmp) CExpression(pmp, GPOS_NEW(pmp) CPatternTree(pmp))
));
```



# Pre-Condition Check



```
// Compatibility function for splitting aggregates
virtual
BOOL FCompatible(CXform::EXformId exfid)
{
    return (CXform::ExfSplitGbAgg != exfid);
}
```



Do not fire this rule on a logical operator produced by the same rule.  
(Avoid Infinite Recursion)

# The Actual Transformation



```
void Transform
(
    CXformContext *pxfctxt, // update
    CXformResult *pxfres, // output
    CExpression *pexpr // input
)
const;
```

details: libgpopt/src/xforms/CXformSplitGbAgg.cpp



# Register Transformation Rule



```
void CXformFactory::Instantiate()  
{  
    ...  
    Add(GPOS_NEW(m_pmp) CXformSplitGbAgg(m_pmp));  
    ...  
}
```



GPORCA



**BUILD**



# Dependencies

GP-XERCES:

<https://github.com/greenplum-db/gp-xerces>

CMake 3.0+



# CMake Build

```
mkdir build
```

```
cd build
```

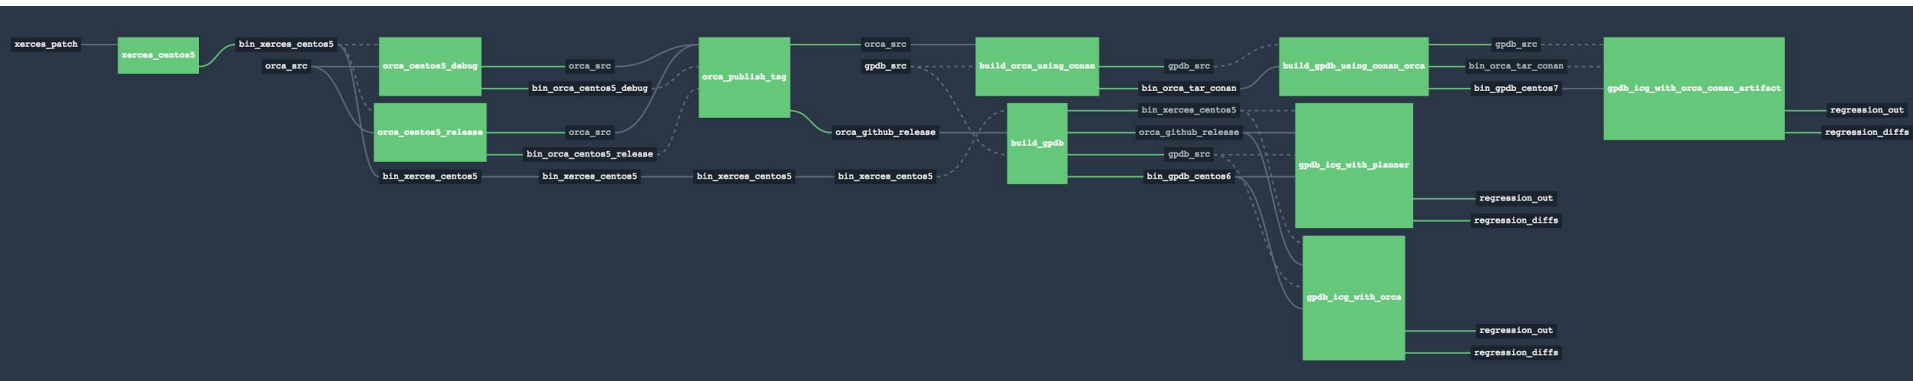
```
cmake ../
```

```
make && make install
```



# CI: Concourse Pipeline

<https://ci.orca.pivotalci.info/teams/main/pipelines/gporca>





GPORCA



**TEST**



# Test GPORCA (<1min)

```
# run all unit tests  
ctest
```

```
# run all unit tests in parallel with 7 threads  
ctest -j7
```

```
# run only one unit test called CAggTest  
./server/gporca_test -U CAggTest
```



# Test with GPDB OSS in Docker

Follow instructions from:

<https://github.com/d/bug-free-fortnight>

It's very useful to verify installcheck-good locally with latest GPORCA changes.





GPORCA

---

**CONTRIBUTE**



## 1-2-3

1. Fork GPORCA at

<https://github.com/greenplum-db/gporca>

2. Pick an issue

<https://github.com/greenplum-db/gporca/issues>

3. Send a Pull Request (PR)

A group of people in a workshop setting. A man on the left is pointing at a wall covered in sticky notes. A group of four people (two men and two women) are seated in the center, looking towards the left. A man on the right is standing and looking towards the group. The background shows a workshop environment with various tools and materials.

**Thanks to all these contributors**

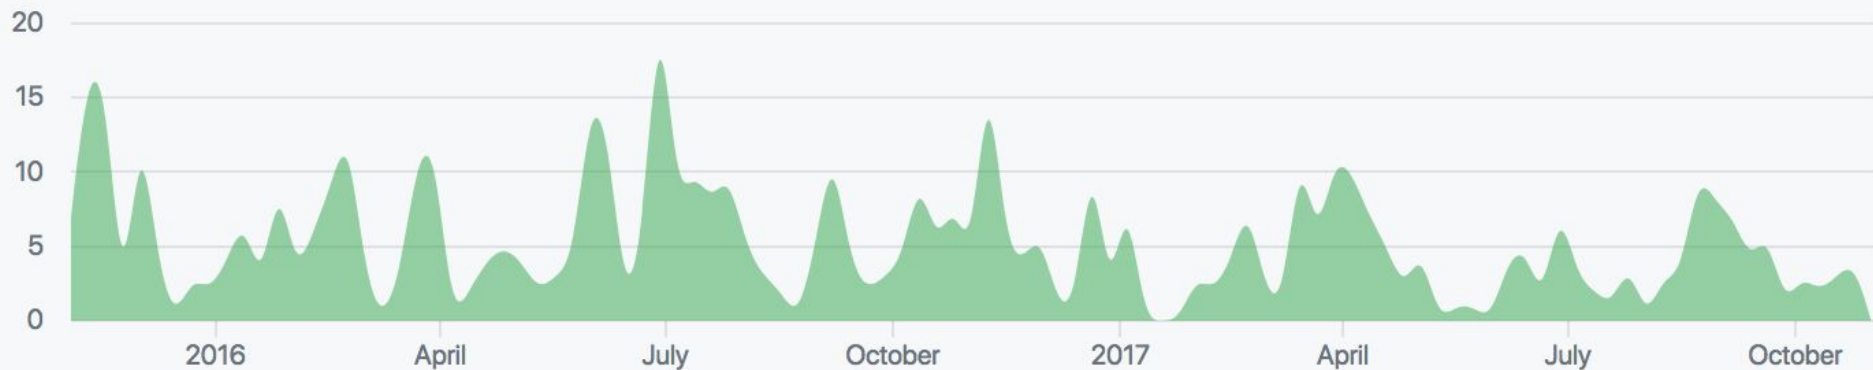


# Mike Waas

Founder & CEO  
Datometry

---

<http://www.ceotodaymagazine.com/issue/issue-09-2017/#32>



d

82 commits 12,282 ++ 7,067 --

#1



oarap

79 commits 217,312 ++ 210,374 --

#2



hsyuan

59 commits 31,536 ++ 30,571 --

#3



craig-chasseur

53 commits 1,949 ++ 1,202 --

#4



vraghavan78

51 commits 18,104 ++ 118,940 --

#5



xinzweb

35 commits 12,151 ++ 21,663 --

#6



**hlinnaka**

#7

32 commits 3,209 ++ 11,699 --



**bhuvnesh2703**

#8

26 commits 9,749 ++ 6,233 --



**dhanashreek89**

#9

25 commits 26,879 ++ 19,223 --



**hardikar**

#10

17 commits 5,880 ++ 665 --



**karthijrk**

#11

15 commits 1,211 ++ 1,389 --



**zaksoup**

#12

13 commits 493 ++ 171 --



**cramja**

#13

11 commits 1,118 ++ 361 --



**khannaekta**

#14

8 commits 1,557 ++ 1,505 --



**jpatel-pivotal**

#15

8 commits 6,269 ++ 107 --



**ryantang**

#16

7 commits 46 ++ 7 --



**lpetrov-pivotal**

7 commits 978 ++ 284 --

#17



**cjcjameson**

6 commits 589 ++ 80 --

#18



**armenatzoglou**

4 commits 170 ++ 135 --

#19



**atris**

3 commits 124 ++ 117 --

#20



**sambitesh**

3 commits 728 ++ 5 --

#21



**addisonhuddy**

2 commits 34 ++ 31 --

#22



**danielgustafsson**

1 commit 2 ++ 2 --

#23



**challiwill**

1 commit 4 ++ 4 --

#24



A group of people in a workshop setting. A man on the left is pointing at a wall covered in sticky notes. Several people are seated in the center, and two more are standing on the right. The scene is dimly lit with a blue tint.

# You are invited!





# THANK YOU

---

GPDB:

<http://greenplum.org/>

<https://github.com/greenplum-db/gpdb>

GPORCA Github:

<https://github.com/greenplum-db/gporca>

mailing lists (400+ member):

[gpdb-users@greenplum.org](mailto:gpdb-users@greenplum.org)

Greenplum YouTube:

<https://www.youtube.com/GreenplumDatabase>

The background of the slide is a teal-colored image of the Golden Gate Bridge, viewed from a low angle looking up at one of the towers. The bridge's cables and structure are visible, extending into the distance.

# Pivotal®



Transforming How The World Builds Software



**Q: What happen to SORT generated in Enforcement?**

A: There is only ONE implementation of SORT, so, that's a physical operator and won't be pushed anywhere after enforcement.

**Q: Why CXformResult has more than one alternatives?**

A: Some rule (e.g. CXformExpandNAryJoinDP) can produce multiple choices after transformation

**Q: How hard to make GPORCA adapt to a new host?**

A: The hard part is on the MD translation. So far, GPORCA is very PostgreSQL friendly.

**Q: Why there is no separate SQL parser included in GPORCA?**

A: GPORCA focused on relational algebra and let host handle the binding, view expansions, and permissions.

**Q: How to add a new property like 'reliability' or 'dollar cost' to this optimizer?**

A: That can be added to Property Enforcement as the Order/Distribution/Partition/Rewindability. For example, if people want to favor a more 'reliable' data source, they can add a CEnfdReliability class to cost that choice. It's an interesting combination of 'reliability' and 'dollar cost', usually, when it's more reliable is more expensive. It's an interesting balance to achieve.

**Q: How long does the `ctest` run?**

A: Around 5min on 2.8Ghz Intel i7. Running with `ctest -j7` finished in < 2min.

**Q: Is GPORCA multi-threaded?**

A: It's multi-thread READY, but currently, we run with single thread. There are still few caveats (thread safe issues) to iron out before we can fully turn it on.



## **Orca: A Modular Query Optimizer Architecture for Big Data, SIGMOD 2014**

Mohamed A. Soliman, Lyublena Antova, Venkatesh Raghavan, Amr El-Helw, Zhongxian Gu, Entong Shen, George C. Caragea, Carlos Garcia-Alvarado, Foyzur Rahman, Michalis Petropoulos, Florian Waas, Sivaramakrishnan Narayanan, Konstantinos Krikellas, Rhonda Baldwin

## **Optimization of Common Table Expressions in MPP Database Systems, VLDB 2015**

Amr El-Helw, Venkatesh Raghavan, Mohamed A. Soliman, George C. Caragea, Zhongxian Gu, Michalis Petropoulos.

## **Optimizing Queries over Partitioned Tables in MPP Systems, SIGMOD 2014**

Lyublena Antova, Amr El-Helw, Mohamed Soliman, Zhongxian Gu, Michalis Petropoulos, Florian Waas

## **Reversing Statistics for Scalable Test Databases Generation, DBTest 2013**

Entong Shen, Lyublena Antova

## **Total Operator State Recall - Cost-Effective Reuse of Results in Greenplum Database, ICDE Workshops 2013**

George C. Caragea, Carlos Garcia-Alvarado, Michalis Petropoulos, Florian M. Waas

## **Testing the Accuracy of Query Optimizers, DBTest 2012**

Zhongxian Gu, Mohamed A. Soliman, Florian M. Waas

## **Automatic Capture of Minimal, Portable, and Executable Bug Repros using AMPERe, DBTest 2012**

Lyublena Antova, Konstantinos Krikellas, Florian M. Waas

## **Automatic Data Placement in MPP Databases, ICDE Workshops 2012**

Carlos Garcia-Alvarado, Venkatesh Raghavan, Sivaramakrishnan Narayanan, Florian M. Waas



GPORCA



**DEBUG**



# GPORCA Traceflags and GPDB GUC

GPORCA relies on Traceflags to change runtime behavior: **Traceflag.h**

Exposed in GPDB as GUC (Grand Unified Configuration): **guc\_gp.c**

```
-- turn on GPORCA  
  
set optimizer=on;  
  
-- print input query (GPORCA TF 101000)  
  
set optimizer_print_query=on;
```



# Minidump: DXL document

Turn on the minidump

```
set client_min_messages='log';  
set optimizer=on;  
set optimizer_enable_constant_expression_evaluation=off;  
set optimizer_enumerate_plans=on;  
set optimizer_minidump=always;
```

Run a query

GPORCA creates a \*.mdp file in the **\$MASTER\_DATA\_DIRECTORY/minidump**

```
# run only one minidump directly  
./server/gporca_test -d ../data/dxl/minidump/TVFRandom.mdp
```



# Input Query, Output Plan

Debug the plans

```
set client_min_messages='log';
```

```
set optimizer=on;
```

```
set optimizer_print_query=on; -- input query, and  
preprocessed query
```

```
set optimizer_print_plan=on; -- output final physical plan
```





# Plan Enumeration

Turn on the plan enumerations

```
set client_min_messages='log';  
set optimizer=on;  
set optimizer_enumerate_plans=on;
```

Pick a plan out of search space

```
set optimizer=on;  
set client_min_messages='log';  
set optimizer_enumerate_plans=on;  
set optimizer_plan_id=1;
```



# Optimization Stats and Xform Rules

Debug optimizer stages

```
set client_min_messages='log';  
set optimizer=on;  
set optimizer_print_optimization_stats=on;
```

Debug the transformation rules details

```
set client_min_messages='log';  
set optimizer=on;  
set optimizer_print_xform=on;
```



# MEMO Groups

```
set optimizer_print_memo_after_exploration=on;  
set optimizer_print_memo_after_implementation=on;  
set optimizer_print_memo_after_optimization=on;
```

ROOT group is indicated as `ROOT`



# Way to Disable Xform Rules

```
select disable_xform('CXformJoinAssociativity');  
select enable_xform('CXformJoinAssociativity');
```

All the xform rules can be found from the class names under  
**libgpopt/include/gpopt/xforms**

**CXformFactory::Instantiate** lists all the activated xform rules (~130 rules)



# Useful Breakpoints

```
# Entry point of optimizer
    COptimizer::PdxlnOptimize
# DXL: Translate DXL into Query
    CTranslatorDXLToExpr::PexprTranslateQuery
# Step 1: Pre-processor
    CExpressionPreprocessor::PexprPreprocess
# Step 2-3-4: Optimization
    COptimizer::PexprOptimize
# Individual rule transformation, all CXform* classes
    CXformSplitGbAgg::Transform
# Enforceable Property
    CEngine::FCheckEnfdProps
    CPartitionPropagationSpec::AppendEnforcers
# DXL: Translate Plan back in DXL
    CTranslatorExprToDXL::PdxlnTranslate
```



GPORCA



# CODE BASE

# Top Level



```
.  
├── cmake  
├── concourse  
├── data  
├── libgpdbcost  
├── libgpopt  
├── libgpos  
├── libnaucrates  
├── patches  
├── scripts  
└── server
```



# libgpos: memory management, task scheduler, exception handling, unit-test framework

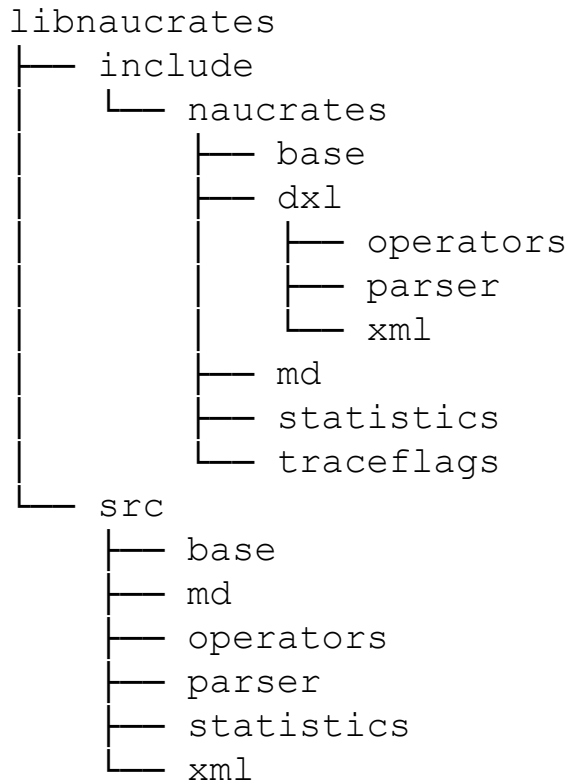
```
libgpos
├── include
└── src
    ├── common
    ├── error
    ├── io
    ├── memory
    ├── net
    ├── string
    ├── sync
    ├── task
    └── test
```

```
server
├── include
└── src
    ├── startup
    └── unittest
        └── gpos
            ├── common
            ├── error
            ├── io
            ├── memory
            ├── string
            ├── sync
            ├── task
            └── test
```





# libnaucrates: DXL, metadata, statistics, traceflags





# libgpopt: engine, metadata cache, minidump, operators, memo, xform rules

```
libgpopt
├── include
│   └── gpopt
└── src
    ├── base
    ├── engine
    ├── eval
    ├── mdcache
    ├── metadata
    ├── minidump
    ├── operators
    ├── optimizer
    ├── search
    ├── translate
    └── xforms
```



# libgpdbcost: cost model

```
libgpdbcost
├── CMakeLists.txt
├── include
│   └── gpdbcost
└── src
    ├── CCostModelGPDB.cpp
    ├── CCostModelGPDBLegacy.cpp
    ├── CCostModelParamsGPDB.cpp
    ├── CCostModelParamsGPDBLegacy.cpp
    └── ICostModel.cpp
```



# server: unit tests

```
server
├── include
├── src
│   ├── startup
│   └── unittest
│       ├── dxl
│       │   ├── base
│       │   └── statistics
│       └── gpopt
│           ├── base
│           ├── cost
│           ├── csq
│           ├── engine
│           ├── eval
│           ├── mdcache
│           ├── metadata
│           ├── minidump
│           ├── operators
│           ├── search
│           ├── translate
│           └── xforms
```



# Data: all the test data

```
data
├── dxl
│   ├── cost
│   ├── csq_tests
│   ├── expressiontests
│   ├── indexjoin
│   ├── metadata
│   ├── minidump
│   │   ├── CArrayExpansionTest
│   │   ├── CJoinOrderDPTTest
│   │   ├── CPhysicalParallelUnionAllTest
│   │   ├── CPruneColumnsTest
│   │   └── sql
│   ├── multilevel-partitioning
│   ├── parse_tests
│   ├── plstmt
│   ├── query
│   ├── search
│   ├── statistics
│   ├── tpcds
│   ├── tpcds-partitioned
│   ├── tpch
│   └── tpch-partitioned
```